

Binary provenance of the shipped wheels (cuobjdump, no GPU required)
torch 2.9.0+cu128 / CUDA 12.8 / flash-attn 2.8.3 / flashinfer 0.6.18 -- an orange cell runs another architecture's cubin under
CUDA minor-version binary compatibility; this is a packaging fact, consistent with (not proof of) the latency it sits beside

		target compute capability				
		sm_80	sm_86	sm_89	sm_90	sm_120
cublas		SASS	SASS	sm_86	SASS	SASS
cudnn-engines		SASS	SASS	sm_86	SASS	SASS
flash-attn-2		SASS	sm_80	sm_80	SASS	SASS
flash-attn-3		n/a	n/a	n/a	n/a	n/a
flashinfer (JIT)		JIT	JIT	JIT	JIT	fail
libtorch_cuda		SASS	SASS	SASS	SASS	SASS

- native SASS
- runs older cubin
- PTX only (JIT at load)
- no code for this arch
- JIT-compiled at run time
- JIT fails (nvcc < 12.9)
- library not installed